Deep Generative Models: Discrete Latent Variables

Philip Schulz and Wilker Aziz

VI Tutorial @ Host
Site

1. First Attempt: Wake-Sleep

2. Neural Variational Inference

3. Score function estimator

4. Variance reduction

Generative Models

Joint distribution over observed data x and latent variables z.

$$p(x, z | \theta) = \underbrace{p(z)}_{\text{prior}} \underbrace{p(x | z, \theta)}_{\text{likelihood}}$$

The likelihood and prior are often standard distributions (Gaussian, Bernoulli) with simple dependence on conditioning information.

Deep generative models

Joint distribution with **deep observation model**

$$p(x, z|\theta) = \underbrace{p(z)}_{\text{prior}} \underbrace{p(x|z, \theta)}_{\text{likelihood}}$$

mapping from z to $p(x|z, \theta)$ is a NN with parameters θ

Deep generative models

Joint distribution with **deep observation model**

$$p(x, z | \theta) = \underbrace{p(z)}_{\text{prior}} \underbrace{p(x | z, \theta)}_{\text{likelihood}}$$

mapping from z to $p(x | z, \theta)$ is a NN with parameters θ

Marginal likelihood

$$p(x | \theta) = \int p(x, z | \theta) \, \mathrm{d}z = \int p(z) p(x | z, \theta) \, \mathrm{d}z$$

intractable in general

Goals

We want
- richer probabilistic models

Goals

We want

- richer probabilistic models
- complex observation models parameterised by NNs

Goals

We want

- richer probabilistic models
- complex observation models parameterised by NNs

but we can't perform gradient-based MLE

Goals

We want

- richer probabilistic models
- complex observation models parameterised by NNs

but we can't perform gradient-based MLE

We need approximate inference techniques!

ELBO recap

And we've developed the ELBO

$$\log p(x) = \log \int p(z, x) \mathrm{d}z$$

ELBO recap

And we've developed the ELBO

$$\log p(x) = \log \int p(z,x)\mathrm{d}z = \log \int q(z|x)\frac{p(z,x)}{q(z|x)}\mathrm{d}z$$

ELBO recap

And we've developed the ELBO

$$\log p(x) = \log \int p(z, x)\mathrm{d}z = \log \int q(z|x)\frac{p(z, x)}{q(z|x)}\mathrm{d}z$$

$$\overset{\mathsf{JI}}{\geq} \underbrace{\mathbb{E}_{q(z|x)}\left[\log \frac{p(z, x)}{q(z|x)}\right]}_{\text{ELBO}}$$

ELBO recap

And we've developed the ELBO

$$\log p(x) = \log \int p(z, x)\mathrm{d}z \; = \log \int q(z|x)\frac{p(z, x)}{q(z|x)}\mathrm{d}z$$

$$\overset{\mathrm{JI}}{\geq} \underbrace{\mathbb{E}_{q(z|x)}\left[\log \frac{p(z, x)}{q(z|x)}\right]}_{\text{ELBO}}$$

$$\mathrm{ELBO} = \mathbb{E}_{q(z|x)}\left[\log \frac{p(z|x)p(x)}{q(z|x)}\right]$$

$$= -\underbrace{\mathrm{KL}\left(q(z|x) \;||\; p(z|x)\right)}_{\text{gap}} + \log p(x)$$

1. **First Attempt: Wake-Sleep**

2. Neural Variational Inference

3. Score function estimator

4. Variance reduction

Wake-sleep Algorithm

- Generalise latent variables to Neural Networks
- Train generative neural model
- Use variational inference! (kind of)

Wake-sleep Architecture

2 Neural Networks:

Wake-sleep Architecture

2 Neural Networks:

- A generation network to model the data (the one we want to optimise) – parameters: θ

Wake-sleep Architecture

2 Neural Networks:

- A generation network to model the data (the one we want to optimise) – parameters: θ
- An inference (recognition) network (to model the latent variable) – parameters: λ

Wake-sleep Architecture

2 Neural Networks:

- A generation network to model the data (the one we want to optimise) – parameters: θ
- An inference (recognition) network (to model the latent variable) – parameters: λ
- Original setting: binary hidden units

Wake-sleep Architecture

2 Neural Networks:

- A generation network to model the data (the one we want to optimise) – parameters: θ
- An inference (recognition) network (to model the latent variable) – parameters: λ
- Original setting: binary hidden units
- Training is performed in a "hard EM" fashion

Generator

Inference Network

Wake-sleep Training

Wake Phase

- Use inference network to sample hidden unit setting z from $q(z|x, \lambda)$
- Update generation parameters θ to maximize joint log-liklelihood of data and latents $p(x, z|\theta)$

Wake-sleep Training

Wake Phase

- Use inference network to sample hidden unit setting z from $q(z|x, \lambda)$
- Update generation parameters θ to maximize joint log-liklelihood of data and latents $p(x, z|\theta)$

Sleep Phase

- Produce dream sample $\tilde{x}$ from random hidden unit z
- Update inference parameters λ to maximize probability of latent state $q(z|\tilde{x}, \lambda)$

Wake Phase Objective

Objective

$$\arg\max_{\theta} \; \mathbb{E}_{p(x)}\left[\text{ELBO}(\theta, \lambda | x)\right]$$

$$= \arg\max_{\theta} \; \mathbb{E}_{p(x)}\left[\mathbb{E}_{q(z|x,\lambda)}\left[\log p(z, x|\theta)\right] + \mathbb{H}[q(z|x,\lambda)]\right]$$

Wake Phase Objective

Objective

$$\arg\max_\theta \; \mathbb{E}_{p(x)} \left[\text{ELBO}(\theta, \lambda | x) \right]$$

$$= \arg\max_\theta \; \mathbb{E}_{p(x)} \left[\mathbb{E}_{q(z|x,\lambda)} \left[\log p(z, x | \theta) \right] + \mathbb{H}[q(z|x,\lambda)] \right]$$

Gradient wrt θ for $x \sim p(x)$

$$\nabla_\theta \mathbb{E}_{q(z|x,\lambda)} \left[\log p(z, x | \theta) \right] + \nabla_\theta \mathbb{H}[q(z|x,\lambda)]$$

Wake Phase Objective

Objective

$$\arg \max_{\theta} \; \mathbb{E}_{p(x)} \left[\text{ELBO}(\theta, \lambda | x) \right]$$

$$= \arg \max_{\theta} \; \mathbb{E}_{p(x)} \left[\mathbb{E}_{q(z|x,\lambda)} \left[\log p(z, x | \theta) \right] + \mathbb{H}[q(z|x,\lambda)] \right]$$

Gradient wrt θ for $x \sim p(x)$

$$\boldsymbol{\nabla}_{\theta} \mathbb{E}_{q(z|x,\lambda)} \left[\log p(z, x | \theta) \right] + \boldsymbol{\nabla}_{\theta} \mathbb{H}[q(z|x,\lambda)]$$

$$= \mathbb{E}_{q(z|x,\lambda)} \left[\boldsymbol{\nabla}_{\theta} \log p(z, x | \theta) \right]$$

Wake Phase Objective

Objective

$$\arg\max_{\theta} \ \mathbb{E}_{p(x)}\left[\text{ELBO}(\theta, \lambda|x)\right]$$

$$= \arg\max_{\theta} \ \mathbb{E}_{p(x)}\left[\mathbb{E}_{q(z|x,\lambda)}\left[\log p(z, x|\theta)\right] + \mathbb{H}[q(z|x,\lambda)]\right]$$

Gradient wrt θ for $x \sim p(x)$

$$\boldsymbol{\nabla}_{\theta}\mathbb{E}_{q(z|x,\lambda)}\left[\log p(z, x|\theta)\right] + \boldsymbol{\nabla}_{\theta}\mathbb{H}[q(z|x,\lambda)]$$

$$= \mathbb{E}_{q(z|x,\lambda)}\left[\boldsymbol{\nabla}_{\theta}\log p(z, x|\theta)\right]$$

$$\overset{\text{MC}}{\approx} \boldsymbol{\nabla}_{\theta}\log p(z, x|\theta) \quad \text{where } z \sim q(z|x,\lambda)$$
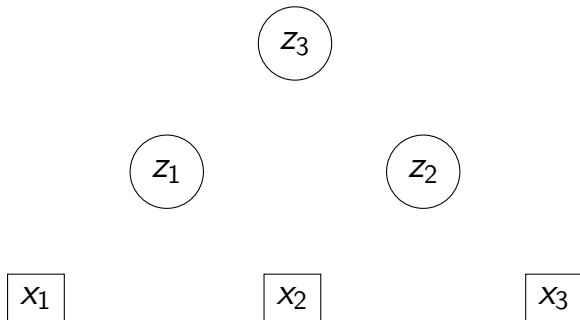
Wake Phase Objective

Assumes z to be fixed random draw from $q(z|x, \lambda)$ and maximises $\log p(z, x|\theta)$.

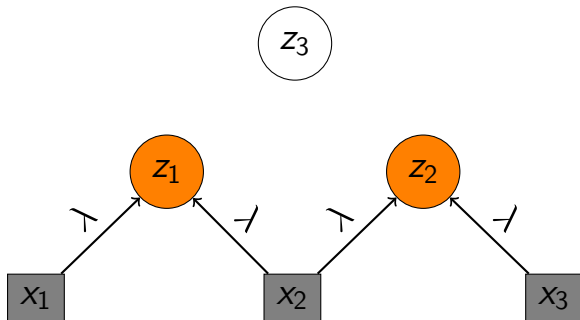This is simply supervised learning with imputed latent data!

Wake Phase Sampling

Sampling $z \sim q(z|x, \lambda)$

Wake Phase Sampling

Sampling $z \sim q(z|x, \lambda)$

Wake Phase Sampling

Sampling $z \sim q(z|x, \lambda)$

Wake Phase Sampling

Sampling $z \sim q(z|x, \lambda)$

Wake Phase Update

Compute $\log p(x, z|\theta)$ and update θ

Sleep Phase Objective

Objective

$$\arg\max_{\lambda}\ \mathbb{E}_{p(x)}\left[\text{ELBO}(\theta, \lambda | x)\right]$$

$$= \arg\max_{\lambda}\ \mathbb{E}_{p(x)}\left[\mathbb{E}_{q(z|x,\lambda)}\left[\log p(z, x|\theta)\right] + \mathbb{H}[q(z|x,\lambda)]\right]$$

Sleep Phase Objective

Objective

$$\arg\max_{\lambda} \; \mathbb{E}_{p(x)}\left[\text{ELBO}(\theta, \lambda | x)\right]$$

$$= \arg\max_{\lambda} \; \mathbb{E}_{p(x)}\left[\mathbb{E}_{q(z|x,\lambda)}\left[\log p(z, x | \theta)\right] + \mathbb{H}[q(z|x,\lambda)]\right]$$

Gradient wrt λ for $x \sim p(x)$

$$\boldsymbol{\nabla}_{\lambda}\mathbb{E}_{q(z|x,\lambda)}\left[\log p(z, x | \theta)\right] + \boldsymbol{\nabla}_{\lambda}\mathbb{H}[q(z|x,\lambda)]$$

Sleep Phase Objective

Objective

$$\arg\max_{\lambda} \; \mathbb{E}_{p(x)}\left[\mathsf{ELBO}(\theta, \lambda | x)\right]$$

$$= \arg\max_{\lambda} \; \mathbb{E}_{p(x)}\left[\mathbb{E}_{q(z|x,\lambda)}\left[\log p(z, x | \theta)\right] + \mathbb{H}[q(z|x,\lambda)]\right]$$

Gradient wrt λ for $x \sim p(x)$

$$\boldsymbol{\nabla}_{\lambda}\mathbb{E}_{q(z|x,\lambda)}\left[\log p(z, x | \theta)\right] + \boldsymbol{\nabla}_{\lambda}\mathbb{H}[q(z|x,\lambda)]$$

Let's change the objective!

Sleep Phase (Convenient) Objective

Flip the direction of the KL

$$\arg\min_{\lambda} \mathbb{E}_{p(x)} \left[\text{KL} \left(p(z|x,\theta) \, || \, q(z|x,\lambda) \right) \right]$$

Sleep Phase (Convenient) Objective

Flip the direction of the KL

$$\arg\min_{\lambda} \mathbb{E}_{p(x)} \left[\mathrm{KL}\left(p(z|x,\theta) \,||\, q(z|x,\lambda) \right) \right]$$

$$= \arg\min_{\lambda} \mathbb{E}_{p(x)} \mathbb{E}_{p(z|x,\theta)} \left[\log p(z|x,\theta) - \log q(z|x,\lambda) \right]$$

Sleep Phase (Convenient) Objective

Flip the direction of the KL

$$\arg\min_{\lambda} \mathbb{E}_{p(x)} \left[\text{KL} \left(p(z|x, \theta) \ || \ q(z|x, \lambda) \right) \right]$$

$$= \arg\min_{\lambda} \mathbb{E}_{p(x)} \mathbb{E}_{p(z|x,\theta)} \left[\log p(z|x, \theta) - \log q(z|x, \lambda) \right]$$

$$= \arg\max_{\lambda} \mathbb{E}_{p(x,z|\theta)} \left[\log q(z|x, \lambda) \right] - \underbrace{\mathbb{E}_{p(x,z|\theta)} \left[\log p(z|x, \theta) \right]}_{\text{constant}}$$

Sleep Phase (Convenient) Objective

Flip the direction of the KL

$$\arg\min_{\lambda} \mathbb{E}_{p(x)} \left[\text{KL} \left(p(z|x,\theta) \parallel q(z|x,\lambda) \right) \right]$$

$$= \arg\min_{\lambda} \mathbb{E}_{p(x)} \mathbb{E}_{p(z|x,\theta)} \left[\log p(z|x,\theta) - \log q(z|x,\lambda) \right]$$

$$= \arg\max_{\lambda} \mathbb{E}_{p(x,z|\theta)} \left[\log q(z|x,\lambda) \right] - \underbrace{\mathbb{E}_{p(x,z|\theta)} \left[\log p(z|x,\theta) \right]}_{\text{constant}}$$

Gradient wrt λ

$$\nabla_{\lambda} \mathbb{E}_{p(x,z|\theta)} \left[\log q(z|x,\lambda) \right]$$

Sleep Phase (Convenient) Objective

Flip the direction of the KL

$$\arg\min_{\lambda} \mathbb{E}_{p(x)} \left[\mathrm{KL} \left(p(z|x,\theta) \,||\, q(z|x,\lambda) \right) \right]$$

$$= \arg\min_{\lambda} \mathbb{E}_{p(x)} \mathbb{E}_{p(z|x,\theta)} \left[\log p(z|x,\theta) - \log q(z|x,\lambda) \right]$$

$$= \arg\max_{\lambda} \mathbb{E}_{p(x,z|\theta)} \left[\log q(z|x,\lambda) \right] - \underbrace{\mathbb{E}_{p(x,z|\theta)} \left[\log p(z|x,\theta) \right]}_{\text{constant}}$$

Gradient wrt λ

$$\nabla_{\lambda} \mathbb{E}_{p(x,z|\theta)} \left[\log q(z|x,\lambda) \right]$$

$$= \mathbb{E}_{p(x,z|\theta)} \left[\nabla_{\lambda} \log q(z|x,\lambda) \right]$$

Sleep Phase (Convenient) Objective

Flip the direction of the KL

$$\arg\min_{\lambda} \mathbb{E}_{p(x)}\left[\text{KL}\left(p(z|x,\theta) \,||\, q(z|x,\lambda)\right)\right]$$

$$= \arg\min_{\lambda} \mathbb{E}_{p(x)}\mathbb{E}_{p(z|x,\theta)}\left[\log p(z|x,\theta) - \log q(z|x,\lambda)\right]$$

$$= \arg\max_{\lambda} \mathbb{E}_{p(x,z|\theta)}\left[\log q(z|x,\lambda)\right] - \underbrace{\mathbb{E}_{p(x,z|\theta)}\left[\log p(z|x,\theta)\right]}_{\text{constant}}$$

Gradient wrt λ

$$\nabla_{\lambda}\mathbb{E}_{p(x,z|\theta)}\left[\log q(z|x,\lambda)\right]$$

$$= \mathbb{E}_{p(x,z|\theta)}\left[\nabla_{\lambda}\log q(z|x,\lambda)\right]$$

$$\overset{\text{MC}}{\approx} \nabla_{\lambda}\log q(z|\tilde{x},\lambda) \quad \text{where } z \sim p(z|\theta)$$

$$\tilde{x} \sim p(x|z,\theta)$$

Sleep Phase (Convenient) Objective

Assumes fake data $\tilde{x}$ and latent variables z to be fixed random draws from $p(x, z|\theta)$ via

$$z \sim p(z|\theta)$$
$$\tilde{x} \sim p(x|z, \theta)$$

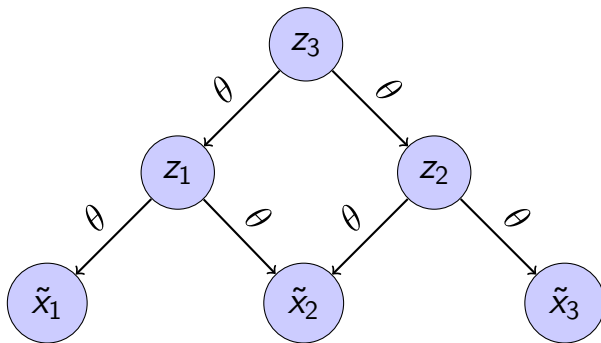and maximises $\log q(z|\tilde{x}, \lambda)$.

Sleep Phase Sampling

Sampling $(z, \tilde{x}) \sim p(x, z | \theta)$

Sleep Phase Sampling

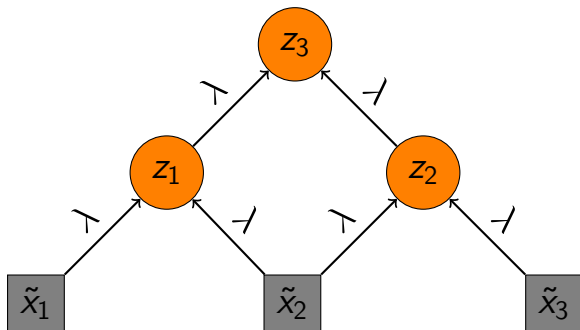Sampling $(z, \tilde{x}) \sim p(x, z|\theta)$

Sleep Phase Sampling

Sampling $(z, \tilde{x}) \sim p(x, z | \theta)$

Sleep Phase Update

Compute $\log q(z|\tilde{x}, \lambda)$ and update λ

Wake-sleep Algorithm

Advantages

- Simple layer-wise updates
- Amortised inference: all latent variables are inferred from the same weights λ

Wake-sleep Algorithm

Advantages

- Simple layer-wise updates
- Amortised inference: all latent variables are inferred from the same weights λ

Drawbacks

- Inference and generative networks are trained on different objectives
- Inference weights λ are updated on fake data $\tilde{x}$
- Generative weights are bad initially, giving wrong signal to the updates of λ

1. First Attempt: Wake-Sleep

2. Neural Variational Inference

3. Score function estimator

4. Variance reduction

Discrete DGMs

Variational Inference Learning (NVIL)

Generative model with NN likelihood

Variational Inference Learning (NVIL)

Generative model with NN likelihood

Let us consider a latent factor model for topic modelling:

Variational Inference Learning (NVIL)

Generative model with NN likelihood

Let us consider a latent factor model for topic modelling:

- a document $x = (x_1, \ldots, x_N)$ consists of n i.i.d. categorical draws from that model

Variational Inference Learning (NVIL)

Generative model with NN likelihood

Let us consider a latent factor model for topic modelling:

- a document $x = (x_1, \ldots, x_N)$ consists of n i.i.d. categorical draws from that model
- the categorical distribution in turn depends on binary latent factors $z = (z_1, \ldots, z_K)$ which are also i.i.d.

Latent factor model

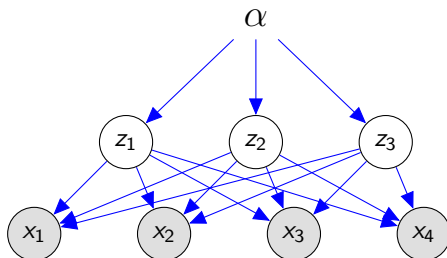$$Z_j \sim \text{Bernoulli}\,(\alpha) \qquad (1 \leq k \leq K)$$
$$X_i | z \sim \text{Categorical}\,(f(z; \theta)) \quad (1 \leq i \leq N)$$

Here $0 < \alpha < 1$ specifies a Bernoulli prior and $f(\cdot; \theta)$ is a function computed by a neural network with softmax output, e.g.

$$f(z; \theta) = \text{softmax}(Wz + b)$$
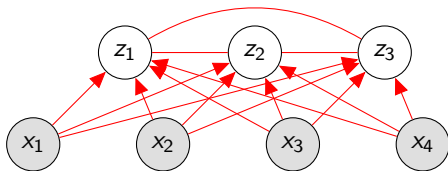$$\theta = \{W, b\}$$

Example Model

Joint distribution: independent latent variables

Example Model

Posterior: latent variables are marginally dependent.

For our variational distribution we are going to assume that they are not (recall: mean field assumption).

Mean Field Inference

The inference network needs to predict K Bernoulli parameters b_1^K. Any neural network with sigmoid output will do that job.

Inference Network

$$q(z|x, \lambda) = \prod_{k=1}^{K} \text{Bern}(z_k | b_k)$$

$$\text{where } b_1^K = g(x; \lambda)$$

Example architecture

$$h = \frac{1}{N} \sum_{i=1}^{N} E_{x_i} \qquad b_1^K = \text{sigmoid}(Mh + c)$$

$$\lambda = \{E, M, c\}$$

Objective

$$\text{ELBO} = \mathbb{E}_{q(z|x,\lambda)}\left[\log p(x,z|\theta)\right] + \mathbb{H}\left(q(z|x,\lambda)\right)$$

Objective

$$\text{ELBO} = \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x, z|\theta) \right] + \mathbb{H} \left(q(z|x, \lambda) \right)$$
$$= \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z, \theta) \right] - \text{KL} \left(q(z|x, \lambda) \, || \, p(z) \right)$$

Objective

$$\text{ELBO} = \mathbb{E}_{q(z|x,\lambda)}\left[\log p(x, z|\theta)\right] + \mathbb{H}\left(q(z|x, \lambda)\right)$$
$$= \mathbb{E}_{q(z|x,\lambda)}\left[\log p(x|z, \theta)\right] - \text{KL}\left(q(z|x, \lambda) \mid\mid p(z)\right)$$

Parameter estimation

$$\underset{\theta,\lambda}{\arg\max} \ \mathbb{E}_{q(z|x,\lambda)}\left[\log p(x|z, \theta)\right] - \text{KL}\left(q(z|x, \lambda) \mid\mid p(z)\right)$$

KL

KL between K independent Bernoulli distributions is tractable

$$\mathsf{KL}\left(q(z|x,\lambda) \mid\mid p(z|\alpha)\right) = \sum_{k=1}^{K} \mathsf{KL}\left(q(z_k|x,\lambda) \mid\mid p(z_k|\alpha)\right)$$

KL

KL between K independent Bernoulli distributions is tractable

$$\mathrm{KL}\left(q(z|x,\lambda) \;||\; p(z|\alpha)\right) = \sum_{k=1}^{K} \mathrm{KL}\left(q(z_k|x,\lambda) \;||\; p(z_k|\alpha)\right)$$

$$= \sum_{k=1}^{K} \mathrm{KL}\left(\mathrm{Bernoulli}\left(b_k\right)\right) \;||\; \mathrm{Bernoulli}\left(\alpha\right))$$

KL

KL between K independent Bernoulli distributions is tractable

$$\text{KL}\left(q(z|x,\lambda) \,||\, p(z|\alpha)\right) = \sum_{k=1}^{K} \text{KL}\left(q(z_k|x,\lambda) \,||\, p(z_k|\alpha)\right)$$

$$= \sum_{k=1}^{K} \text{KL}\left(\text{Bernoulli}\left(b_k\right)\right) \,||\, \text{Bernoulli}\left(\alpha\right))$$

$$= \sum_{k=1}^{K} b_k \log \frac{b_k}{\alpha} + (1 - b_k) \log \frac{1 - b_k}{1 - \alpha}$$

Generative Network Gradient

$$\frac{\partial}{\partial \theta} \left(\mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta) \right] - \overbrace{\mathsf{KL}\left(q(z|x,\lambda) \,||\, p(z) \right)}^{\text{constant wrt } \theta} \right)$$

Generative Network Gradient

$$\frac{\partial}{\partial \theta} \left(\mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta) \right] - \overbrace{\mathsf{KL}\left(q(z|x,\lambda) \mid\mid p(z) \right)}^{\text{constant wrt } \theta} \right)$$

$$= \underbrace{\mathbb{E}_{q(z|x,\lambda)} \left[\frac{\partial}{\partial \theta} \log p(x|z,\theta) \right]}_{\text{expected gradient :)}}$$

Generative Network Gradient

$$\frac{\partial}{\partial \theta} \left(\mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta)\right] - \overbrace{\text{KL}\left(q(z|x,\lambda) \,\|\, p(z)\right)}^{\text{constant wrt } \theta} \right)$$

$$= \underbrace{\mathbb{E}_{q(z|x,\lambda)} \left[\frac{\partial}{\partial \theta} \log p(x|z,\theta)\right]}_{\text{expected gradient :)}}$$

$$\overset{\text{MC}}{\approx} \frac{1}{S} \sum_{s=1}^{S} \frac{\partial}{\partial \theta} \log p(x|z^{(s)},\theta) \quad \text{where } z^{(s)} \sim q(z|x,\lambda)$$

Inference Network Gradient

$$\frac{\partial}{\partial \lambda} \left(\mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta) \right] - \overbrace{\text{KL}\left(q(z|x,\lambda) \,||\, p(z) \right)}^{\text{analytical}} \right)$$

Inference Network Gradient

$$\frac{\partial}{\partial \lambda} \left(\mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta) \right] - \overbrace{\mathsf{KL}\left(q(z|x,\lambda) \mid\mid p(z)\right)}^{\text{analytical}} \right)$$

$$= \frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta) \right] - \underbrace{\frac{\partial}{\partial \lambda} \mathsf{KL}\left(q(z|x,\lambda) \mid\mid p(z)\right)}_{\text{analytical computation}}$$

Inference Network Gradient

$$\frac{\partial}{\partial \lambda} \left(\mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta) \right] - \overbrace{\text{KL} \left(q(z|x,\lambda) \mid\mid p(z) \right)}^{\text{analytical}} \right)$$

$$= \frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta) \right] - \underbrace{\frac{\partial}{\partial \lambda} \text{KL} \left(q(z|x,\lambda) \mid\mid p(z) \right)}_{\text{analytical computation}}$$

The first term again requires approximation by sampling, but there is a problem

Inference Network Gradient

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q_\lambda(z|x)} \left[\log p_\theta(x|z) \right]$$

Inference Network Gradient

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q_\lambda(z|x)} \left[\log p_\theta(x|z) \right]$$

$$= \frac{\partial}{\partial \lambda} \sum_z q(z|x, \lambda) \log p(x|z, \theta)$$

Inference Network Gradient

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q_\lambda(z|x)} \left[\log p_\theta(x|z) \right]$$

$$= \frac{\partial}{\partial \lambda} \sum_z q(z|x, \lambda) \log p(x|z, \theta)$$

$$= \underbrace{\sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta)}_{\text{not an expectation}}$$

Inference Network Gradient

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q_\lambda(z|x)} \left[\log p_\theta(x|z) \right]$$

$$= \frac{\partial}{\partial \lambda} \sum_z q(z|x, \lambda) \log p(x|z, \theta)$$

$$= \underbrace{\sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta)}_{\text{not an expectation}}$$

- MC estimator is non-differentiable

Inference Network Gradient

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q_\lambda(z|x)} \left[\log p_\theta(x|z) \right]$$

$$= \frac{\partial}{\partial \lambda} \sum_z q(z|x, \lambda) \log p(x|z, \theta)$$

$$= \underbrace{\sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta)}_{\text{not an expectation}}$$

- MC estimator is non-differentiable
- Differentiating the expression does not yield an expectation: cannot approximate via MC

Score function estimator

We can again use the log identity for derivatives

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q_\lambda(z|x)} \left[\log p_\theta(x|z) \right]$$

$$= \sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta)$$

Score function estimator

We can again use the log identity for derivatives

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q_\lambda(z|x)} \left[\log p_\theta(x|z) \right]$$

$$= \sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta)$$

$$= \sum_z q(z|x, \lambda) \frac{\partial}{\partial \lambda} (\log q(z|x, \lambda)) \log p(x|z, \theta)$$

Score function estimator

We can again use the log identity for derivatives

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q_\lambda(z|x)} \left[\log p_\theta(x|z) \right]$$

$$= \sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta)$$

$$= \sum_z q(z|x, \lambda) \frac{\partial}{\partial \lambda} (\log q(z|x, \lambda)) \log p(x|z, \theta)$$

$$= \underbrace{\mathbb{E}_{q(z|x, \lambda)} \left[\log p(x|z, \theta) \frac{\partial}{\partial \lambda} \log q(z|x, \lambda) \right]}_{\text{expected gradient :)}}$$

Score function estimator: remarks

We can now build an MC estimator

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta) \right]$$

$$= \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta) \frac{\partial}{\partial \lambda} \log q(z|x,\lambda) \right]$$

Score function estimator: remarks

We can now build an MC estimator

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z, \theta) \right]$$

$$= \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z, \theta) \frac{\partial}{\partial \lambda} \log q(z|x, \lambda) \right]$$

$$\overset{\text{MC}}{\approx} \frac{1}{S} \sum_{s=1}^{S} \log p(x|z^{(s)}, \theta) \frac{\partial}{\partial \lambda} \log q(z^{(s)}|x, \lambda)$$

where $z^{(s)} \sim q(z|x, \lambda)$

Score function estimator: remarks

We can now build an MC estimator

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta) \right]$$

$$= \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z,\theta) \frac{\partial}{\partial \lambda} \log q(z|x,\lambda) \right]$$

$$\overset{\text{MC}}{\approx} \frac{1}{S} \sum_{s=1}^{S} \log p(x|z^{(s)},\theta) \frac{\partial}{\partial \lambda} \log q(z^{(s)}|x,\lambda)$$

where $z^{(s)} \sim q(z|x,\lambda)$

- magnitude of $\log p(x|z,\theta)$ varies widely
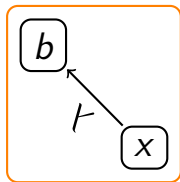
Score function estimator: remarks

We can now build an MC estimator

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z, \theta) \right]$$

$$= \mathbb{E}_{q(z|x,\lambda)} \left[\log p(x|z, \theta) \frac{\partial}{\partial \lambda} \log q(z|x, \lambda) \right]$$

$$\overset{\text{MC}}{\approx} \frac{1}{S} \sum_{s=1}^{S} \log p(x|z^{(s)}, \theta) \frac{\partial}{\partial \lambda} \log q(z^{(s)}|x, \lambda)$$
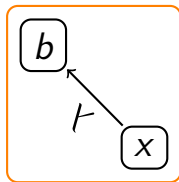
where $z^{(s)} \sim q(z|x, \lambda)$

- magnitude of $\log p(x|z, \theta)$ varies widely
- model likelihood does not contribute to direction of gradient
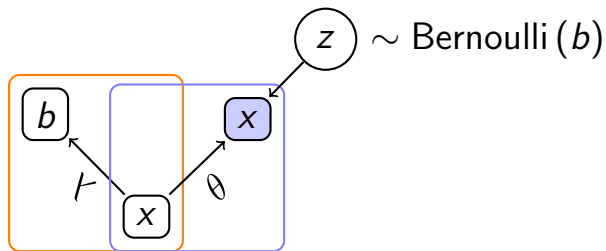
Computation Graph

inference model

Computation Graph

$$z \sim \text{Bernoulli}\,(b)$$

inference model
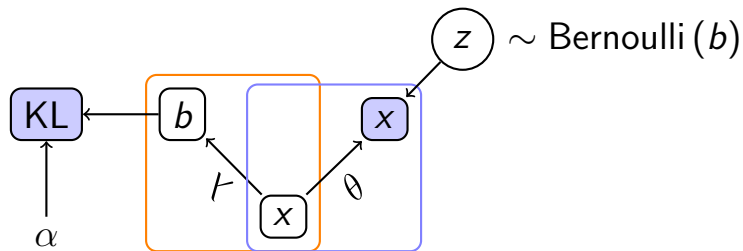
Computation Graph

inference model   generation model

Computation Graph
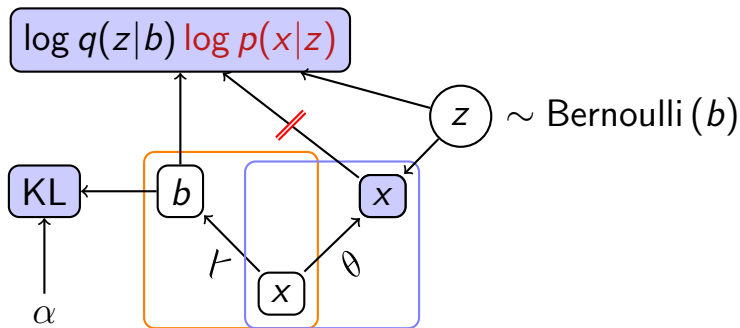

inference model   generation model

Computation Graph

inference model   generation model

Pros and Cons

- Pros
  - Applicable to all distributions
  - Many libraries come with samplers for common distributions

Pros and Cons

- Pros
  - Applicable to all distributions
  - Many libraries come with samplers for common distributions
- Cons
  - High Variance!

1. First Attempt: Wake-Sleep

2. Neural Variational Inference

3. Score function estimator

4. Variance reduction

When variance is high we can

When variance is high we can

- sample more

When variance is high we can

- sample more
- use variance reduction techniques (e.g. baselines and control variates)

Control variates

Intuition

To estimate $\mathbb{E}[f(z)]$ via Monte Carlo we compute the empirical average of $\hat{f}(z)$ where $\hat{f}(z)$ is chosen so that $\mathbb{E}[\hat{f}(z)] = \mathbb{E}[f(z)]$ and $\text{Var}(f) > \text{Var}(\hat{f})$.

Equivalent expectations

Let $\bar{f} = \mathbb{E}[f(z)]$ be an expectation of interest

Equivalent expectations

Let $\bar{f} = \mathbb{E}[f(z)]$ be an expectation of interest

- say we know $\bar{c} = \mathbb{E}[c(z)]$

Equivalent expectations

Let $\bar{f} = \mathbb{E}[f(z)]$ be an expectation of interest

- say we know $\bar{c} = \mathbb{E}[c(z)]$
- then for $\hat{f}(z) \triangleq f(z) - b(c(z) - \mathbb{E}[c(z)])$

Equivalent expectations

Let $\bar{f} = \mathbb{E}[f(z)]$ be an expectation of interest

- say we know $\bar{c} = \mathbb{E}[c(z)]$
- then for $\hat{f}(z) \triangleq f(z) - b(c(z) - \mathbb{E}[c(z)])$
  it holds that $\mathbb{E}[\hat{f}(z)] = \mathbb{E}[f(z)]$

Equivalent expectations

Let $\bar{f} = \mathbb{E}[f(z)]$ be an expectation of interest

- say we know $\bar{c} = \mathbb{E}[c(z)]$
- then for $\hat{f}(z) \triangleq f(z) - b(c(z) - \mathbb{E}[c(z)])$
  it holds that $\mathbb{E}[\hat{f}(z)] = \mathbb{E}[f(z)]$
- and $\text{Var}(\hat{f}) = \text{Var}(f) - 2b\,\text{Cov}(f, c) + b^2\,\text{Var}(c)$

Choosing the control variate

1. $\hat{f}(z) \triangleq f(z) - b(c(z) - \mathbb{E}[c(z)])$
2. $\text{Var}(\hat{f}) = \text{Var}(f) - 2b\,\text{Cov}(f, c) + b^2\,\text{Var}(c)$

How do we choose b and $c(z)$?

Choosing the control variate

1. $\hat{f}(z) \triangleq f(z) - b(c(z) - \mathbb{E}[c(z)])$
2. $\text{Var}(\hat{f}) = \text{Var}(f) - 2b\,\text{Cov}(f, c) + b^2\,\text{Var}(c)$

How do we choose b and $c(z)$?

- If $f(z)$ and $c(z)$ are positively correlated, then we may reduce variance

Choosing the control variate

1. $\hat{f}(z) \triangleq f(z) - b(c(z) - \mathbb{E}[c(z)])$
2. $\text{Var}(\hat{f}) = \text{Var}(f) - 2b\,\text{Cov}(f, c) + b^2\,\text{Var}(c)$

How do we choose b and $c(z)$?

- If $f(z)$ and $c(z)$ are positively correlated, then we may reduce variance
- solving $\frac{\partial}{\partial b}\,\text{Var}(\hat{f}) = 0$

Choosing the control variate

1. $\hat{f}(z) \triangleq f(z) - b(c(z) - \mathbb{E}[c(z)])$
2. $\text{Var}(\hat{f}) = \text{Var}(f) - 2b\,\text{Cov}(f, c) + b^2\,\text{Var}(c)$

How do we choose b and $c(z)$?

- If $f(z)$ and $c(z)$ are positively correlated, then we may reduce variance
- solving $\frac{\partial}{\partial b}\text{Var}(\hat{f}) = 0$ yields $b^\star = {}^{\text{Cov}(f, c)}/_{\text{Var}(c)}$

Choosing the control variate

1. $\hat{f}(z) \triangleq f(z) - b(c(z) - \mathbb{E}[c(z)])$
2. $\mathsf{Var}(\hat{f}) = \mathsf{Var}(f) - 2b\,\mathsf{Cov}(f, c) + b^2\,\mathsf{Var}(c)$

How do we choose b and $c(z)$?

- If $f(z)$ and $c(z)$ are positively correlated, then we may reduce variance
- solving $\frac{\partial}{\partial b}\mathsf{Var}(\hat{f}) = 0$ yields $b^\star = {}^{\mathsf{Cov}(f,\,c)}\!/_{\mathsf{Var}(c)}$

Of course, $\mathbb{E}[c(z)]$ must be known!

MC

We then use the estimate

$$\bar{f} \overset{\text{MC}}{\approx} \frac{1}{S} \left(\sum_{s=1}^{S} f(z^{(s)}) - bc(z^{(s)}) \right) + b\bar{c}$$

MC

We then use the estimate

$$\bar{f} \overset{\text{MC}}{\approx} \frac{1}{S} \left(\sum_{s=1}^{S} f(z^{(s)}) - bc(z^{(s)}) \right) + b\bar{c}$$

And recall that for us

$$f(z) = \log p(x|z, \theta) \frac{\partial}{\partial \lambda} \log q(z|x, \lambda)$$

and $z^{(s)} \sim q(z|x, \lambda)$

Expected score

The Expectation of the score function is 0.

$$\mathbb{E}_{q(z|x,\lambda)} \left[\frac{\partial}{\partial \lambda} \log q(z|x,\lambda) \right]$$

Expected score

The Expectation of the score function is 0.

$$\mathbb{E}_{q(z|x,\lambda)} \left[\frac{\partial}{\partial \lambda} \log q(z|x,\lambda) \right]$$

$$= \int q(z|x,\lambda) \frac{\partial}{\partial \lambda} \log q(z|x,\lambda) \mathrm{d}z$$

Expected score

The Expectation of the score function is 0.

$$\mathbb{E}_{q(z|x,\lambda)} \left[\frac{\partial}{\partial \lambda} \log q(z|x,\lambda) \right]$$

$$= \int q(z|x,\lambda) \frac{\partial}{\partial \lambda} \log q(z|x,\lambda) \mathrm{d}z$$

$$= \int \frac{\partial}{\partial \lambda} q(z|x,\lambda) \mathrm{d}z$$

Expected score

The Expectation of the score function is 0.

$$\mathbb{E}_{q(z|x,\lambda)}\left[\frac{\partial}{\partial\lambda}\log q(z|x,\lambda)\right]$$

$$= \int q(z|x,\lambda)\frac{\partial}{\partial\lambda}\log q(z|x,\lambda)\mathrm{d}z$$

$$= \int \frac{\partial}{\partial\lambda}q(z|x,\lambda)\mathrm{d}z$$

$$= \frac{\partial}{\partial\lambda}\int q(z|x,\lambda)\mathrm{d}z$$

Expected score

The Expectation of the score function is 0.

$$\mathbb{E}_{q(z|x,\lambda)} \left[\frac{\partial}{\partial \lambda} \log q(z|x,\lambda) \right]$$

$$= \int q(z|x,\lambda) \frac{\partial}{\partial \lambda} \log q(z|x,\lambda) \mathrm{d}z$$

$$= \int \frac{\partial}{\partial \lambda} q(z|x,\lambda) \mathrm{d}z$$

$$= \frac{\partial}{\partial \lambda} \int q(z|x,\lambda) \mathrm{d}z$$

$$= \frac{\partial}{\partial \lambda} 1 = 0$$

Baselines

With

$$f(z) = \log p(x|z,\theta) \frac{\partial}{\partial \lambda} \log q(z|x,\lambda)$$

and

$$c(z) = \frac{\partial}{\partial \lambda} \log q(z|x,\lambda)$$

we have

$$\hat{f}(z) =$$

Baselines

With

$$f(z) = \log p(x|z, \theta) \frac{\partial}{\partial \lambda} \log q(z|x, \lambda)$$

and

$$c(z) = \frac{\partial}{\partial \lambda} \log q(z|x, \lambda)$$

we have

$$\hat{f}(z) = (\log p(x|z, \theta) - b) \frac{\partial}{\partial \lambda} \log q(z|x, \lambda)$$

Baselines

With

$$f(z) = \log p(x|z,\theta)\frac{\partial}{\partial \lambda} \log q(z|x,\lambda)$$

and

$$c(z) = \frac{\partial}{\partial \lambda} \log q(z|x,\lambda)$$

we have

$$\hat{f}(z) = (\log p(x|z,\theta) - b)\frac{\partial}{\partial \lambda} \log q(z|x,\lambda)$$

b is known as *baseline* in RL literature.

Examples of baselines

- Moving average of $\log p(x|z, \theta)$ based on previous batches

Examples of baselines

- Moving average of $\log p(x|z, \theta)$ based on previous batches
- A trainable constant b

Examples of baselines

- Moving average of $\log p(x|z, \theta)$ based on previous batches
- A trainable constant b
- A neural network prediction based on x e.g. $b(x; \omega)$

Examples of baselines

- Moving average of $\log p(x|z, \theta)$ based on previous batches
- A trainable constant b
- A neural network prediction based on x e.g. $b(x; \omega)$
- The likelihood assessed at a deterministic point, e.g. $b(x) = \log p(x|z^\star, \theta)$ where $z^\star = \arg\max_z q(z|x, \lambda)$

Trainable baselines

Baselines are predicted by a regression model (e.g. a neural net).

The model is trained using an L_2-loss.

$$\min_{\omega} \left(b(x; \omega) - \log p(x|z, \theta) \right)^2$$

Summary

- Wake-Sleep: train inference and generation networks with separate objectives

Summary

- Wake-Sleep: train inference and generation networks with separate objectives
- NVIL: a single objective (ELBO) for both models

Summary

- Wake-Sleep: train inference and generation networks with separate objectives
- NVIL: a single objective (ELBO) for both models
- Use score function estimator

Summary

- Wake-Sleep: train inference and generation networks with separate objectives
- NVIL: a single objective (ELBO) for both models
- Use score function estimator
- Always use baselines for variance reduction!

Implementation

Check one of our notebooks, e.g.

- inducing rationales for sentiment classification
  github.com/vitutorial/exercises/tree/
  master/SST

Literature I

G. E. Hinton, P. Dayan, B. J. Frey, and R. M. Neal. The wake-sleep algorithm for unsupervised neural networks. *Science*, 268:1158–1161, 1995. URL http://www.gatsby.ucl.ac.uk/~dayan/papers/hdfn95.pdf.

Andriy Mnih and Karol Gregor. Neural variational inference and learning in belief networks. In *Proceedings of the 31st International Conference on International Conference on Machine Learning - Volume 32*, ICML'14, pages II–1791–II–1799. JMLR.org, 2014. URL http://dl.acm.org/citation.cfm?id=3044805.3045092.

Literature II

John W. Paisley, David M. Blei, and Michael I. Jordan. Variational bayesian inference with stochastic search. In *ICML*, 2012. URL http://icml.cc/2012/papers/687.pdf.

Rajesh Ranganath, Sean Gerrish, and David Blei. Black Box Variational Inference. In Samuel Kaski and Jukka Corander, editors, *AISTATS*, pages 814–822, 2014. URL http://proceedings.mlr.press/v33/ranganath14.pdf.

Literature III

Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8(3-4):229–256, 1992. URL https://doi.org/10.1007/BF00992696.

Evan Greensmith, Peter L Bartlett, and Jonathan Baxter. Variance reduction techniques for gradient estimates in reinforcement learning. *Journal of Machine Learning Research*, 5(Nov):1471–1530, 2004.

Shixiang Gu, Sergey Levine, Ilya Sutskever, and Andriy Mnih. Muprop: Unbiased backpropagation for stochastic neural networks. In *ICLR*, 2016.

Literature IV

George Tucker, Andriy Mnih, Chris J Maddison, John Lawson, and Jascha Sohl-Dickstein. Rebar: Low-variance, unbiased gradient estimates for discrete latent variable models. In *Advances in Neural Information Processing Systems*, pages 2627–2636, 2017.

Will Grathwohl, Dami Choi, Yuhuai Wu, Geoffrey Roeder, and David Duvenaud. Backpropagation through the void: Optimizing control variates for black-box gradient estimation. In *ICLR*, 2018.